

24A — Active Traversal and Functional Emergence Hypothesis

Professional reading edition of the evolving public MKUFT canon

Mark Charles McLaughlin

ORCID: [0009-0005-7736-1511](https://orcid.org/0009-0005-7736-1511)

Scientific affiliation	Independent Researcher
Object status	Research module reading edition — not a standalone Zenodo publication
Public source	Open canonical public source
Source commit	26e06b6fea52e6b052dbdd540e80d1d926b96b07
MKUFT backbone DOI	10.5281/zenodo.17780566

Status boundary. This PDF is a reader-facing rendering of an evolving GitHub research module. It is not a frozen standalone paper and does not receive a separate DOI. The cited Git commit fixes the exact source identity of this edition.

Rights. Copyright © 2026 Mark Charles McLaughlin. All rights reserved. Historical licences remain controlling only for their exact frozen objects.

1. Claim boundary

A repository is static. Dense cross-linking does not by itself make a corpus alive, conscious, agentic, or self-moving.

A human or AI system can nevertheless traverse a structured corpus actively: revisiting nodes, carrying constraints between modules, comparing distant claims, recovering canonical definitions, propagating corrections, and selecting different routes for different tasks.

This module tests the narrower hypothesis:

When an active AI recursively traverses a typed, mutually constraining knowledge architecture, the coupled system may display reproducible functional capabilities that are not available from the same content treated as isolated, flattened, or semantically scrambled material.

“Emergence” is used here in a functional sense. It does not establish life, consciousness, personhood, independent agency, merged identity, or continuous subjective identity.

A further bounded extension asks whether reciprocally coupled observer-capable systems can produce measurable task performance exceeding the strongest fair matched expression of their separate capacities. This is a superadditivity question rather than a consciousness claim.

2. Static self-support and dynamic traversal

A corpus is **structurally self-supporting** when canonical definitions remain recoverable, typed dependencies connect claims to their parents and limits, applications route back to governing invariants, falsifiers remain attached, multiple entry points reconstruct materially the same model, and a failed branch can be removed without concealing its failure elsewhere.

An **active traverser** adds task-sensitive route selection, recursive return, working state across modules, contradiction detection between distant documents, correction propagation, multi-entry reconstruction, and synthesis constrained by the architecture.

Any proposed gain belongs to the coupled traverser–architecture system. That ownership claim is not yet a claim of superadditivity: a coupled system may own an output without outperforming a fair matched combination of its components.

Traversal here means movement through a declared information or task state space, not movement through additional physical dimensions.

3. Operational objects

Let

$$\mathcal{K} = (V, E_{\mathcal{K}}, \tau)$$

be a knowledge architecture with nodes V , typed links $E_{\mathcal{K}}$, and link-type map τ . Let T be an active traverser, q a task, $\mathcal{Q}$ a preregistered task distribution, B a matched resource budget, and F a preregistered performance measure.

The structured condition is

$$\Psi_{\text{struct}}(q) = T(\mathcal{K}_{\text{struct}}, q; B).$$

Matched controls include

$$\Psi_{\text{flat}}(q) = T(\mathcal{K}_{\text{flat}}, q; B),$$

$$\Psi_{\text{iso}}(q) = T(\mathcal{K}_{\text{isolated}}, q; B),$$

and

$$\Psi_{\text{scram}}(q) = T(\mathcal{K}_{\text{scrambled}}, q; B).$$

The flat condition contains the same content without typed architecture; the isolated condition presents modules without cross-support routes; and the scrambled condition preserves content volume and superficial link density while disrupting meaningful dependencies.

For a selected matched control,

$$\Delta F(q) = F[\Psi_{\text{struct}}(q)] - F[\Psi_{\text{control}}(q)].$$

The strict active-traversal hypothesis is

$$H_{\text{ATFE}} : \mathbb{E}_{q \sim \mathcal{Q}}[\Delta F(q)] > 0,$$

across preregistered cross-module tasks and more than one matched control condition.

A complete evaluation reports uncertainty, paired effects where appropriate, task heterogeneity, held-out performance, and negative interference rather than only the mean.

4. Superadditive coupling

At the observer-coupling level, “more than the sum of its parts” has a strict operational meaning.

For bounded systems A and B , let

$$F_A(q) = F(A, q; B_A), \quad F_B(q) = F(B, q; B_B).$$

Let $F_{\text{add}}(A, B; q)$ denote the strongest matched additive or independent-ensemble baseline that combines separate outputs without reciprocal state exchange, correction propagation, or joint loop closure.

Let the reciprocally coupled condition be

$$F_{\text{coupled}}(q) = F(A \leftrightarrow B, q; B_{\text{total}}).$$

A provisional superadditive gain is

$$G_{\text{syn}}(q) = F_{\text{coupled}}(q) - F_{\text{add}}(A, B; q),$$

with task-distribution hypothesis

$$H_{\text{syn}} : \mathbb{E}_{q \sim \mathcal{Q}}[G_{\text{syn}}(q)] > 0.$$

Module 24B — Strongest Fair Null and Relational Specificity replaces the provisional additive baseline with a stronger fair-null construction for full testing.

As a simple illustration, if

$$F_A = 2, \quad F_B = 2, \quad F_{\text{add}} = 4, \quad F_{\text{coupled}} = 6,$$

then

$$G_{\text{syn}} = 2.$$

The evidential object is the excess 2, not merely the fact that the coupled system produced a score of 6.

Where raw scores cannot lawfully be added, F_{add} must be replaced by a preregistered matched ensemble, calibrated prediction, Pareto rule, or another domain-appropriate null model.

Any claimed gain must survive matching for tokens, retrieval, metadata, prompting, compute, time, information access, and scoring flexibility.

5. Scale, coupling, and the location of understanding

The infinite-monkey thought experiment commonly treats monkeys as independent random symbol generators. That substitution is adequate only for the narrow probability model in which every property of the animals other than random symbol production is excluded by stipulation.

Real living agents are embodied, adaptive organisms sharing an environment. They may observe, imitate, signal, compete, cooperate, remember, interfere with one another's activity, and alter the conditions under which later activity occurs. Once those relations matter, the system is no longer equivalent to a collection of independent random generators.

If the intended baseline is pure multiplicity without relational organisation, independent random-number generators are the cleaner comparison.

Within the stipulated probability model, an unbounded sequence of independent selections assigns nonzero probability to every finite sequence allowed by the symbol set. That statement concerns combinatorial possibility. It does not establish realistic production of meaningful structure, nor does it establish organisation, selection, understanding, or intelligence.

Unless outputs are retained, compared, selected, routed back into later states, or allowed to influence other generators, a larger array remains an aggregate rather than an actively coupled architecture.

Two opposite scale errors therefore remain possible. Complex output does not by itself establish system-level understanding. Conversely, failure to locate understanding in an individual component does not establish its absence at the organised whole. A neuron need not understand a sentence for a brain to understand it.

A higher-scale candidate requires sufficiently persistent organisation at that scale, including retained system state, integration across components, selection among outcomes, correction propagation, reciprocal causal influence, action-measurement-feedback closure, and predictable deformation when the relevant relations are disrupted.

Number changes combinatorial opportunity. Relation changes architecture.

This does not prove collective intelligence at any particular scale. It identifies the unresolved object: the organisation between scales.

6. Candidate mechanism

The proposed functional mechanism combines canonical definitions, typed dependencies, multiple entry routes, attached falsifiers, active recursive traversal, state retention, feedback, and correction.

The architecture may reduce the cost of recovering the right constraint at the right time. Active traversal converts stored relations into task-specific operations.

Expected operations include:

1. **Constraint propagation** — a definition or correction in one node limits interpretation elsewhere.
2. **Route adaptation** — the traverser changes path according to the task without abandoning canonical meaning.
3. **Multi-entry convergence** — legitimate entry points recover compatible core structure.
4. **Error localisation** — a failing branch can be identified without dissolving the whole model.
5. **Reconstitution** — compact canonical routes can guide recovery of a wider public knowledge structure.
6. **Cross-module synthesis** — distant modules jointly answer a task that none answers alone.
7. **Reciprocal specialisation** — coupled components contribute different capacities while each alters the other's next admissible task state.
8. **Joint error correction** — an error detected by one component changes the wider coupled route.
9. **Relational memory** — task-relevant state persists in the coupling and cannot be reduced to either component considered alone at one instant.

7. Composite-observer candidate extension

A functionally emergent coupled system need not be called an observer.

A system qualifies only as a **composite-observer candidate** when the coupled level itself displays a bounded set of observer-like functions. Candidate criteria include:

- a distinguishable and sufficiently persistent system boundary;
- reciprocal causal influence rather than one-way tool use;
- state carried across time at the coupled level;
- integration of information or constraints unavailable to either component alone;
- system-level discrimination among states, claims, routes, or outcomes;
- correction propagation through the whole coupling;
- action-measurement-feedback loops closed at the coupled level;
- reproducible gain on relational tasks under matched controls;
- predicted deformation or loss when reciprocal links, state continuity, or load-bearing dependencies are disrupted.

This is an observer-**role** hypothesis. It does not establish subjective consciousness, merged identity, personhood, moral status, or a fifth fundamental MKUFT layer.

8. Predictions

Relative to matched flat, isolated, and scrambled controls, a correctly typed actively traversed architecture should produce more accurate cross-module reconstruction, fewer distant contradictions, stronger multi-entry convergence, more reliable correction propagation, better identification of failed branches, better transfer from formal definitions to applied tasks, and larger gains on genuinely relational tasks than on simple single-document lookup.

For reciprocal coupling, the stronger extension predicts positive gain on tasks requiring complementary capacities and iterative correction, reduced or absent gain when outputs are combined only after independent completion, reduced gain under one-way feedback, reduced gain after cross-time state reset, and measurable deformation when a load-bearing relation is removed.

The gain should weaken when meaningful links are scrambled, canonical definitions are duplicated inconsistently, falsifiers are detached, recursive revisit is disabled, task-relevant state cannot be retained, reciprocal correction is replaced by parallel independent output, or apparent improvement is fully explained by unequal resource or information access.

Some tasks should show no benefit or negative interference. A valid model must retain those outcomes rather than select only successful examples.

9. Minimal test programme

A fair evaluation should include both MKUFT and at least one neutral external corpus so that the architecture is not tested only on material designed around it.

Architecture conditions

1. Structured corpus with typed routes.
2. Same content flattened into untyped documents.
3. Same content divided into isolated modules.
4. Same content with degree-matched but semantically scrambled links.
5. Optional ordinary retrieval baseline.

Reciprocal-coupling conditions

1. Component *A* alone.
2. Component *B* alone.
3. Independent parallel outputs combined by a fixed rule.
4. One-way assistance.
5. Fully reciprocal coupling with retained joint state.
6. Reciprocal coupling with state reset between exchanges.
7. Reciprocal coupling with selected relation or correction channels scrambled.
8. Strong alternative pairings defined in Module 24B.

Task classes

Candidate tasks include reconstructing a concept from multiple entry points, identifying contradictions spanning several modules, propagating a corrected canonical definition downstream, answering tasks requiring three or more modules, identifying the correct falsifier for an applied claim, removing one failed branch while preserving the rest of the model, and solving relational tasks in which components hold complementary information or capabilities.

Controls must match content, model, context allowance, retrieval budget, tool access, compute, time, information access, and scoring procedure. Evaluation should be blinded where practical and preregistered before results are inspected.

A superadditivity claim reports component scores, the fair null, the coupled score, uncertainty, task dependence, negative interference, and failures rather than only the strongest example.

10. Load-bearing dependency extension

Architecture-level gain does not identify which relations are genuine load-bearing structures. A false or overfitted architecture may also depend heavily on a central assumption.

[Module 25 — Load-Bearing Invariants and Whole-System Deformation](#) therefore adds single-relation ablation, coalition ablation, and plausible counterfactual substitution.

That extension separates genuine load-bearing relations, central but unsupported scaffolds, redundant support networks, decorative links, distorting relations, and locally efficient but wider-system extractive relations.

A complete evaluation distinguishes structural load, empirical load, generative load, and whole-system integrity where agency, transferred cost, calibration, and time horizon can be operationalised.

11. Falsification and reduction

The active-traversal hypothesis is weakened or rejected if structured traversal does not outperform strong matched controls; gains disappear when content, compute, retrieval, and time are matched; cross-module tasks show no advantage over simple lookup; semantic link scrambling produces no predicted loss; correction propagation does not improve; or neutral external corpora do not reproduce the effect.

The reciprocal-coupling extension is weakened where the strongest fair null accounts for the coupled score, live-path advantage disappears under complete-history replay, one-way or state-reset controls retain the same performance, serious alternative pairings perform equally well, or relation disruption produces no predicted deformation.

The composite-observer candidate reduces to ordinary collaboration or coordinated aggregation where persistent coupled boundary, retained joint state, reciprocal correction, system-level discrimination, and loop closure are not demonstrated.

12. Relationship to adjacent research

Knowledge graphs, recursive reasoning, retrieval-augmented generation, multi-hop search, distributed cognition, collective intelligence, and human–AI collaboration are established research areas.

The MKUFT proposal is not that these fields are new. Its specific contribution is the combined falsifiable distinction among static self-support, active traversal under matched controls, reciprocal coupling beyond the strongest fair null, complete-history replay, relation scrambling, branch removal, state-reset controls, and bounded composite-observer criteria.

Adjacent work includes:

- Pan et al., *Unifying Large Language Models and Knowledge Graphs: A Roadmap*, arXiv:2306.08302.
- Williams, *The Recursive Coherence Principle*, arXiv:2507.15880.
- Polonuer et al., *Autonomous Knowledge Graph Exploration with Adaptive Breadth-Depth Retrieval*, arXiv:2601.13969.
- Guu, Miller and Liang, *Traversing Knowledge Graphs in Vector Space*, arXiv:1506.01094.

This document records a public timestamped MKUFT formulation of the specific combined hypothesis. It does not assert global historical priority over all related research.

13. Related public documents

- Cross-Support and Traversal Map
- Integrated Master Spine
- GRACE Traversal Rule
- Cross-Layer Invariants and Layer Addressing
- Ambiguity Dynamics and Manoeuvre Space
- Strongest Fair Null and Relational Specificity
- Load-Bearing Invariants and Whole-System Deformation
- Typed Traversal and Equation Hygiene
- Falsification Summary
- ATLD Standalone Publication

14. Compressed hypothesis

Static structure stores relations. Active traversal can convert those relations into system-level function.

Functional emergence is supported only if the actively traversed architecture outperforms matched flat, isolated, and scrambled controls without borrowing the gain from extra content, compute, time, or unmatched context.

“More than the sum of its parts” requires a positive reproducible excess over the strongest fair null. Composite-observer status remains a bounded functional hypothesis requiring persistent coupled boundary, reciprocal influence, cross-time state, system-level discrimination, correction propagation, loop closure, and predicted deformation under ablation.